

PHYS 512: Problem Set 3

Jade Ducharme (260929684)

October 2021

Problem 1

For the first part of this problem, I constructed an RK4 stepper similar to the one we saw in class during Lecture 5. The ODE used to test it is given by Equation 1:

$$\frac{dy}{dx} = \frac{y}{1+x^2} \quad (1)$$

We are given initial condition $y(-20) = 1$. Plugging all this information in *WolframAlpha*, I get that the solution to this ODE (with initial condition) is as shown in Equation 2.

$$y(x) = \exp(\arctan(x) + \arctan(-20)) \quad (2)$$

Figure 1 compares the ODE approximation I get using the RK4 stepper vs. the true solution given by Equation 2:

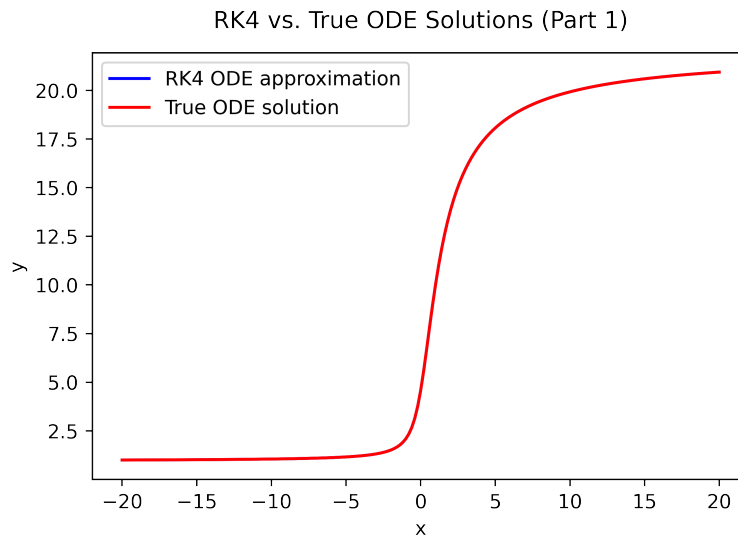


Figure 1: (Blue) approximate solution to the ODE presented in Equation 1 found using an RK4 stepper. (Red) true solution to the same ODE found using *WolframAlpha*. The two curves are indistinguishable to the eye.

Figure 2 shows the residuals for the two curves of Figure 1.

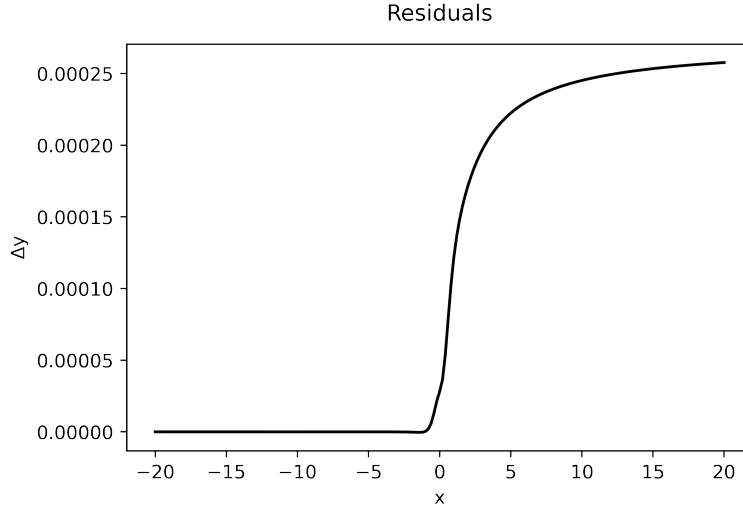


Figure 2: Residuals between the approximate ODE solution and the true ODE solution. The curve tells us the RK4 solution is more accurate for lower values of x .

For the second part of this problem, I need to figure out how to cancel out the leading-order error term from RK4. First, we know the error from a regular RK4 step with step length h is $O(h^5)$. Therefore, if we take an RK4 half-step with length $h_* = \frac{h}{2}$, the error should be $O(h_*^5) = O((\frac{h}{2})^5) = \frac{1}{32}O(h^5)$. But that only takes us from x_n to $x_{n+\frac{1}{2}}$. To get to x_{n+1} , we take an additional half-step, which will also have error $\frac{1}{32}O(h^5)$. So, the total error for both half steps is $2 \times \frac{1}{32}O(h^5) = \frac{1}{16}O(h^5)$.

In order to cancel out the $O(h^5)$ term completely, I need to think of a linear combination of 1 (corresponds to the full step with full error $O(h^5)$) and $\frac{1}{16}$ (corresponds to the two half-steps) that gives zero. The corresponding system of equations is:

$$\begin{cases} a + \frac{b}{16} = 0 \\ a + b = 1 \end{cases} \quad (3)$$

Where the second line comes from the fact that the combination must be normalized, otherwise the solution will blow up. Solving this yields $a = -\frac{1}{15}$ and $b = \frac{16}{15}$. So, if the RK4 approximation from one big step is y_1 , and the approximation from two half-steps is y_2 , the following linear combination cancels out the leading-order error term:

$$y_{n+1} = -\frac{1}{15}y_1 + \frac{16}{15}y_2 \quad (4)$$

Figure 3 shows a plot showing the true vs. approximated ODE curve using Equation 4.

Next, Figure 4 shows the residuals between the true solution and the approximation from Equation 4.

To the eye, there doesn't seem to be a big difference between the approximation using just one full step and the approximation combining two half-steps. In fact, I computed the percent difference between the mean of the residuals for both methods and found that they only differ by about $0.00297 = 0.297\%$. That would imply that using two half-steps instead of one full-step is only about 0.3% more accurate.

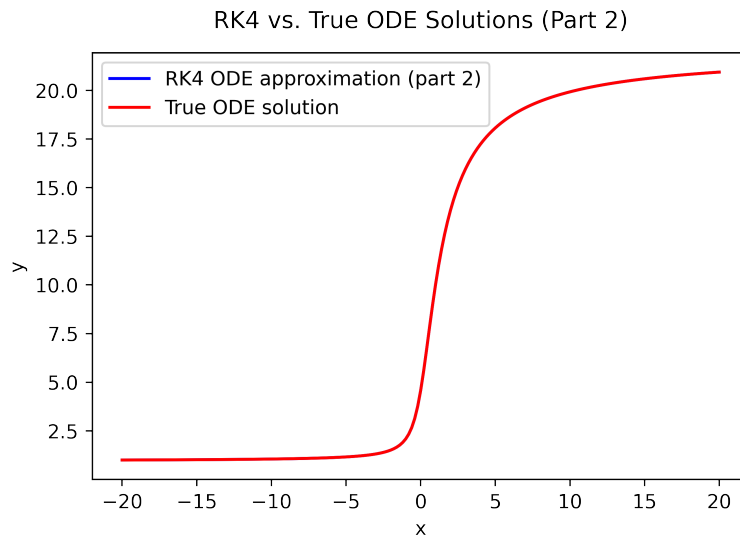


Figure 3: (Blue) approximate solution to the ODE presented in Equation 1 found using Equation 4. (Red) true solution to the same ODE. The two curves are indistinguishable to the eye.




Figure 4: Residuals between the approximate ODE solution and the true ODE solution for Part 2.

The new RK4 stepper now requires 3 function evaluations per step: one to compute the full step with step length h , one to compute a half-step from x_n to $x_{n+\frac{1}{2}}$, and one to compute a half-step from $x_{n+\frac{1}{2}}$ to x_{n+1} . So it's 3 times slower and barely improves the results!

Problem 2

(a) First, I grabbed all the elements of the decay chain as well as their corresponding half-lives and ordered them in a tsv file to keep track of which half-life corresponds to which element. Then, I tweaked the code we saw during Lecture 5 in order to account for every element in the chain. I used `scipy.interpolate.solve_ivp` and specified that I wanted to use the “Radau” method (for implicit integration). That's because I have so

many elements in the decay chain and their decay rates vary so wildly that it would take way too long using *RK4* or other explicit methods.

In Figure 5, I present a few plots.




Figure 5: Rate of change for every element in Uranium-238 decay chain. (Left) The only visible elements are Uranium-238 and Plomb-206. The decay rates of intermediate elements are so short compared to U-238 that they appear to decay instantly and don't show in the plot. (Right) In order to observe the rate of change of the intermediate elements, I place the y-axis on a log scale (technically, a *symlog* scale so that zeros and negative values are treated in a more aesthetically pleasing manner). There are a few negatives near the end of the chain, but I'm pretty sure those are just due to round-off error.

(b) That awful yellow makes it hard to really compare Uranium-238 and Plomb-206, so I present a similar plot with less aggressive colours in Figure 6.

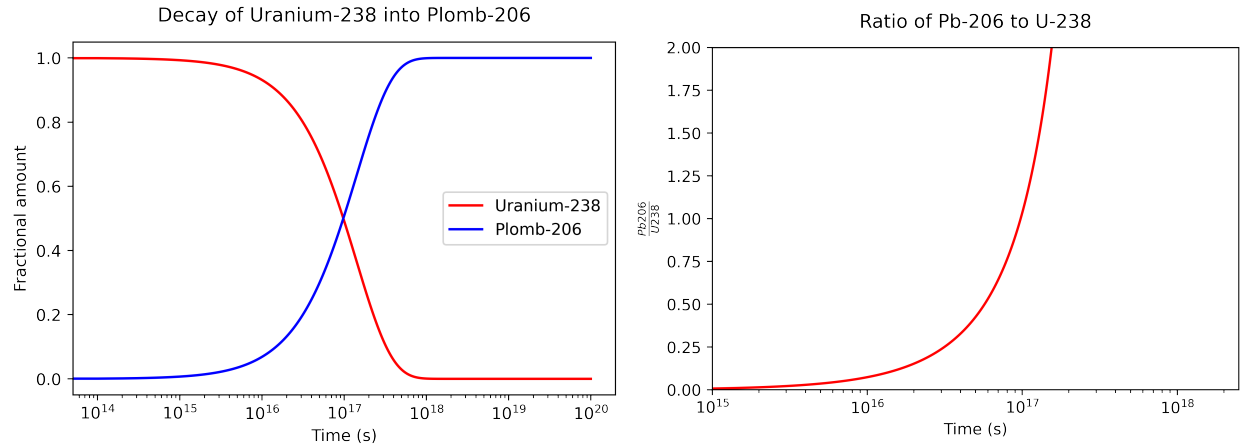


Figure 6: (Left) Decay path for U-238 and Pb-206 (neglecting intermediate elements). (Right) Ratio of Pb-206 to U-238 ($\frac{Pb_{206}}{U_{238}}$).

Analytically, the plots of Figure 6 makes perfect sense: I start off with 100% U-238 and end up with 100% Pb-206. Additionally, the half-life of U-238 is $\sim 10^{17}$ s, which corresponds precisely to the spot where the U-238 curve reaches half-height and the ratio reaches 1. The ratio plot also blows up as the U-238 curve gets closer and closer to zero – also makes sense because U-238 is the denominator. Everything seems to be in order!

Next, I look at the curves for Thorium-230 and Uranium-234 and present them in Figure 7.

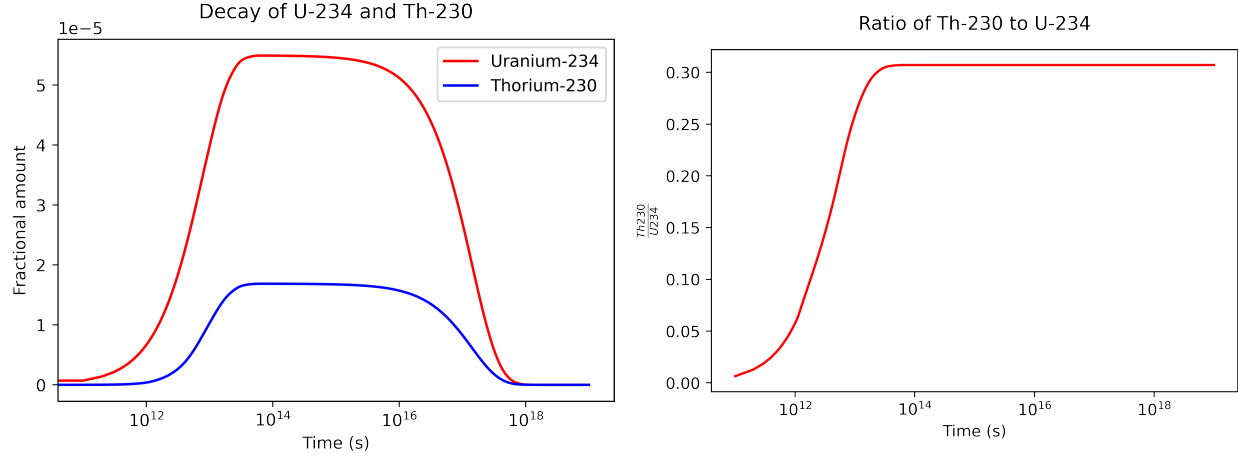


Figure 7: (Left) Decay paths for U-234 and Th-230. (Right) Ratio of Th-230 to U-234 ($\frac{Th^{230}}{U^{234}}$).

We see that the ratio of Th-230 to U-234 reaches equilibrium around the $t = 10^{14}$ s mark, which corresponds to ~ 3.2 million years. So, as long as the sample we're interested in dating isn't older than that, we can use the difference between the ratio of Th-230 to U-234 when the sample was formed and its ratio today to estimate its age.

Problem 3

(a) We have the equation for a symmetric paraboloid:

$$\begin{aligned}
 z - z_0 &= a((x - x_0)^2 + (y - y_0)^2) \\
 z &= z_0 + a(x^2 - 2x_0x + x_0^2) + a(y^2 - 2y_0y + y_0^2) \\
 z &= z_0 + ax^2 - 2ax_0x + ax_0^2 + ay^2 - 2ay_0y + ay_0^2 \\
 z &= (z_0 + ax_0^2 + ay_0^2) + (-2ax_0)x + (-2ay_0)y + a(x^2 + y^2)
 \end{aligned} \tag{5}$$

Therefore, if we define $u \equiv z_0 + ax_0^2 + ay_0^2$, $v \equiv -2ax_0$, and $w \equiv -2ay_0$, the new, *linearized* equation looks like this:

$$z = u + vx + wy + a(x^2 + y^2) \tag{6}$$

We can then obtain an expression for our \mathbf{A} and \mathbf{m} matrices:

$$\mathbf{A} = \begin{pmatrix} 1 & x_1 & y_1 & x_1^2 + y_1^2 \\ 1 & x_2 & y_2 & x_2^2 + y_2^2 \\ 1 & x_3 & y_3 & x_3^2 + y_3^2 \\ \vdots & \vdots & \vdots & \vdots \end{pmatrix} \tag{7}$$

$$\mathbf{m} = \begin{pmatrix} u \\ v \\ w \\ a \end{pmatrix} \quad (8)$$

I know then that the best fit line for my data is given by $z_{fit} = \mathbf{A}\mathbf{m}$. In order to actually *find* the \mathbf{m} matrix, I will use the `numpy.linalg.lstsq` function.

(b) Figure 8 presents plots of z vs. x and y for the raw data as well as the line constructed using $z_{fit} = \mathbf{A}\mathbf{m}$.

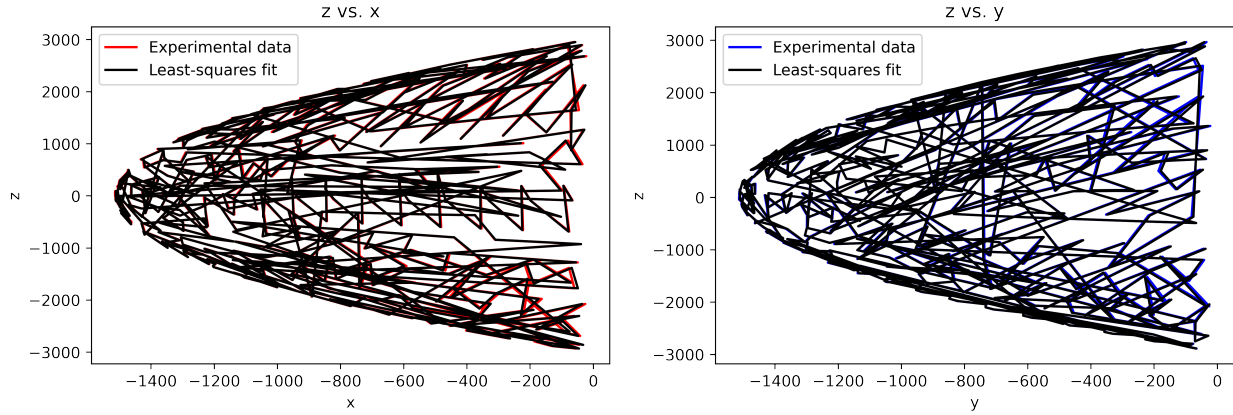


Figure 8: (Left) z vs. x and (right) z vs. y for the raw data (red and blue curves) and the least-squares fit (black curves). To the eye, they look pretty good!

The best fit parameters are:

$$\begin{aligned} u &= -1.51 \times 10^3 \\ v &= 4.54 \times 10^{-4} \\ w &= -1.94 \times 10^{-2} \\ a &= 1.67 \times 10^{-4} \end{aligned} \quad (9)$$

From these, I can easily recover the true parameters for the symmetric parabola:

$$\begin{aligned} a &= 1.67 \times 10^{-4} \\ x_0 &= -1.36 \\ y_0 &= 58.22 \\ z_0 &= -1512.88 \end{aligned} \quad (10)$$

(c) In order to estimate the error on each parameter, we have seen in class that we can define a variable

r as:

$$r = \mathbf{d}_t - \mathbf{A}\mathbf{m}, \quad (11)$$

where \mathbf{d}_t are our true experimental data points and $\mathbf{A}\mathbf{m}$ is the best-fit line given by our model. Here, r simply gives the residuals. We can estimate the errors from the residuals using:

$$\mathbf{N} = \langle r^2 \rangle \quad (12)$$

Ensemble-averaging over r^2 attributes the same error to each data point. This may not be exactly true, but since we have no other information, it's better than nothing. Finally, having obtained our noise matrix \mathbf{N} , the error on the fit parameters is given by $(\mathbf{A}^T \mathbf{N}^{-1} \mathbf{A})^{-1}$, as seen in class.

Using this technique, I found that the uncertainty estimate on the fit parameter a is 6×10^{-8} . This is pretty good!

Next, for the focal point, let's start with our equation for the parabola:

$$\begin{aligned} z &= z_0 + a(x - x_0)^2 + a(y - y_0)^2 \\ z(x) &= a(x - x_0)^2 + \text{constant}, \end{aligned} \quad (13)$$

where in the last line I'm only considering z vs. x . Clearly, neither x_0 nor the $+\text{constant}$ affect the *shape* of the parabola, they simply shift it. Therefore, these constants do not change the focal point. I can then easily rewrite the equation as:

$$z(x) = \frac{x - x_0}{4F} + \text{constant} \quad (14)$$

From this, it becomes evident that $a = \frac{1}{4F}$. Plugging the appropriate value in, I obtain $F = 1499.65$, pretty close to the 1.5m we were hoping for!

Finally, in order to estimate the uncertainty on the focal point, let's say I have a function $F = f(a)$. I will recall the following basic rule for error propagation:

$$\delta F = |f'(a)| \delta a \quad (15)$$

I already know that $F = \frac{1}{4a}$, so the uncertainty on F should be:

$$\begin{aligned} \delta F &= \left| -\frac{1}{4a^2} \right| \delta a \\ F &= \left| -\frac{1}{4(1.67 \times 10^{-4})^2} \right| (6 \times 10^{-8}) \\ F &= 0.6 \end{aligned} \quad (16)$$

Therefore, the value for the focal point is given by $F = 1499.7 \pm 0.6\text{m}$.